

Z-buffer algorithm

for each polygon in model

project vertices of polygon onto viewing plane

for each pixel inside the projected polygon

calculate pixel colour

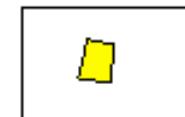
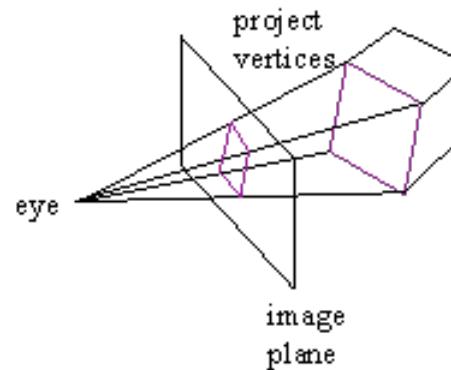
calculate pixel z-value

compare pixel z-value to value stored for pixel in z-buffer

if pixel is closer, draw it in frame-buffer and z-buffer

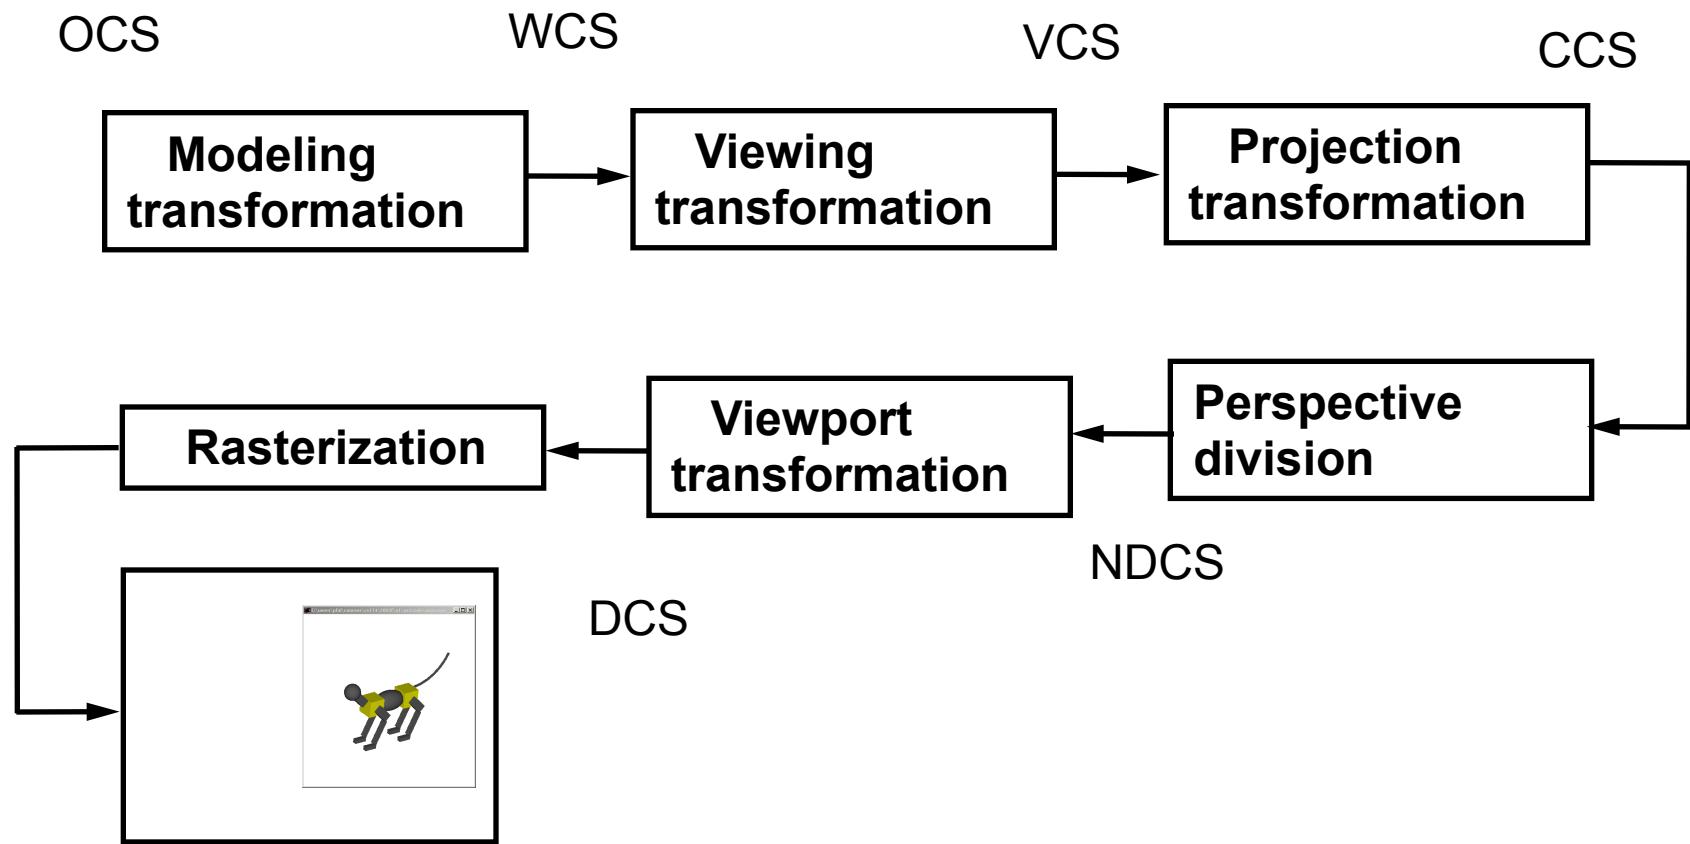
end

end



scan
convert

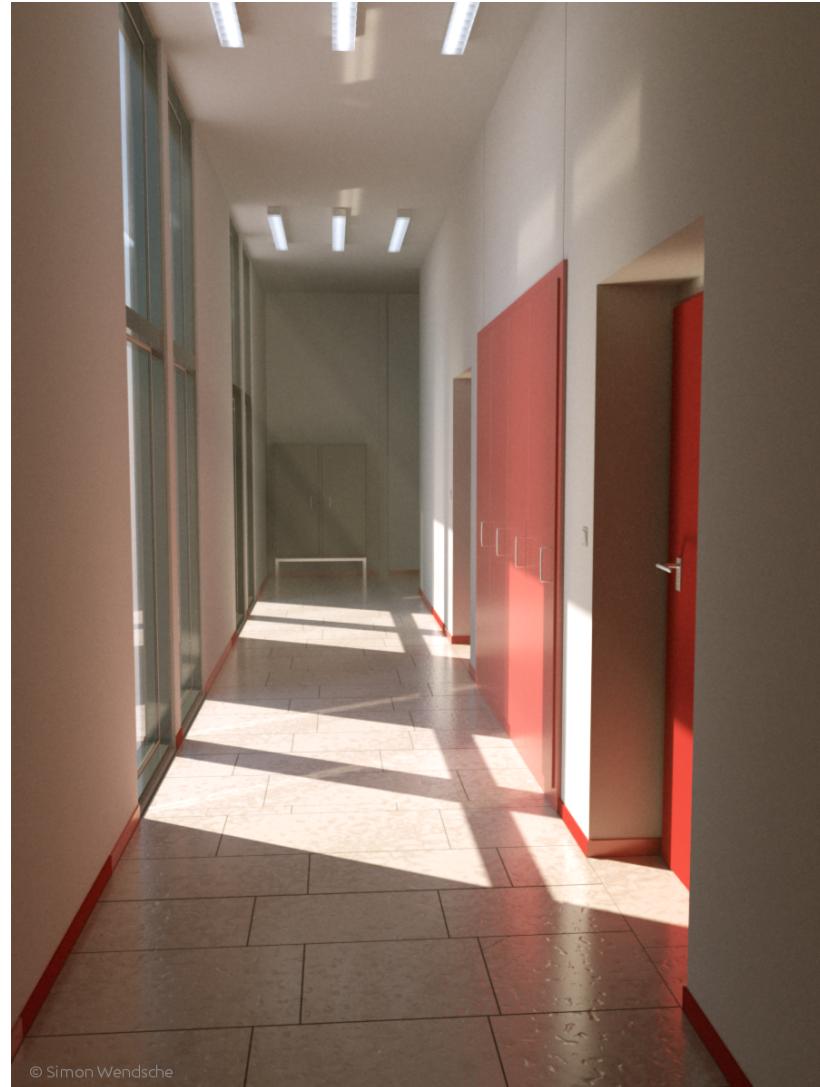
Z-buffer Graphics Pipeline



Global Illumination

Rendered by
LuxRender

*Copyright
Simon Wendsche*



Physics-based Rendering

Light Transport Equation (LTE)

$$L_o(p, \omega_0) = L_e(p, \omega_0) + \int_{\mathcal{S}^2} f(p, \omega_o, \omega_i) L_i(p, \omega_i) |\omega_i \cdot n| d\omega_i$$

L_o : Outgoing radiance at p in direction ω_o

L_e : Emitted radiance at p in direction ω_o

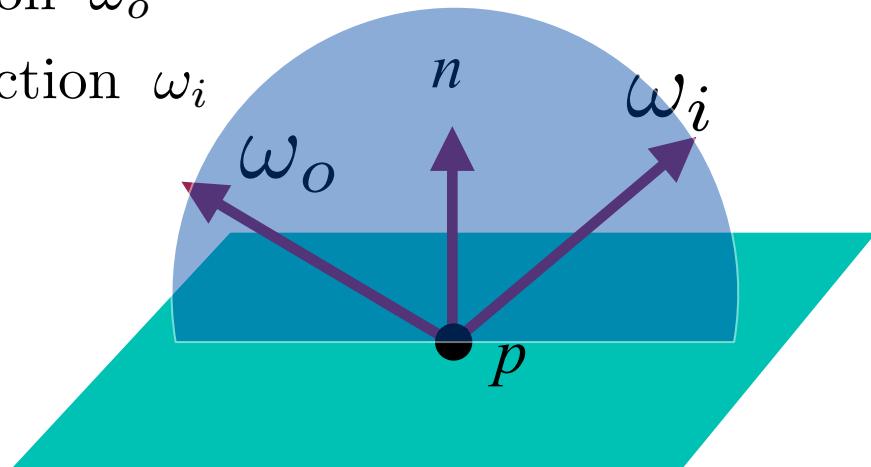
L_i : Incident radiance at p from direction ω_i

n : Unit normal at p

f : Bidirectional scattering function

\mathcal{S}^2 : The unit sphere centered at p

- Integral over directions

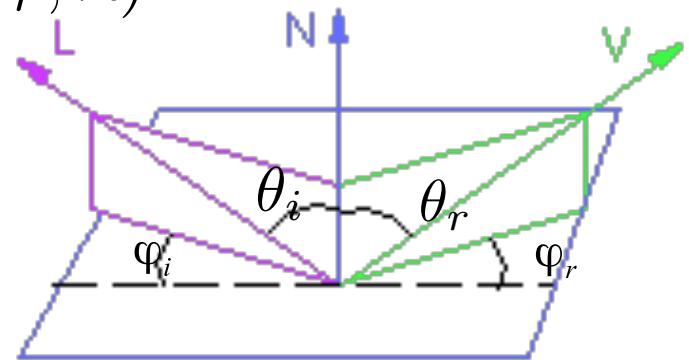


BSDF

Bidirectional Scattering Distribution Function

- BRDF: Bidirectional Reflectance Distribution Function

$$f_r(p, \omega_o, \omega_i, \lambda) = f_r(p, \phi_i, \theta_i, \phi_r, \theta_r, \lambda)$$



- BTDF: Bidirectional Transmittance Distribution function

Examples of a few simple diffuse BRDFs (wikipedia)

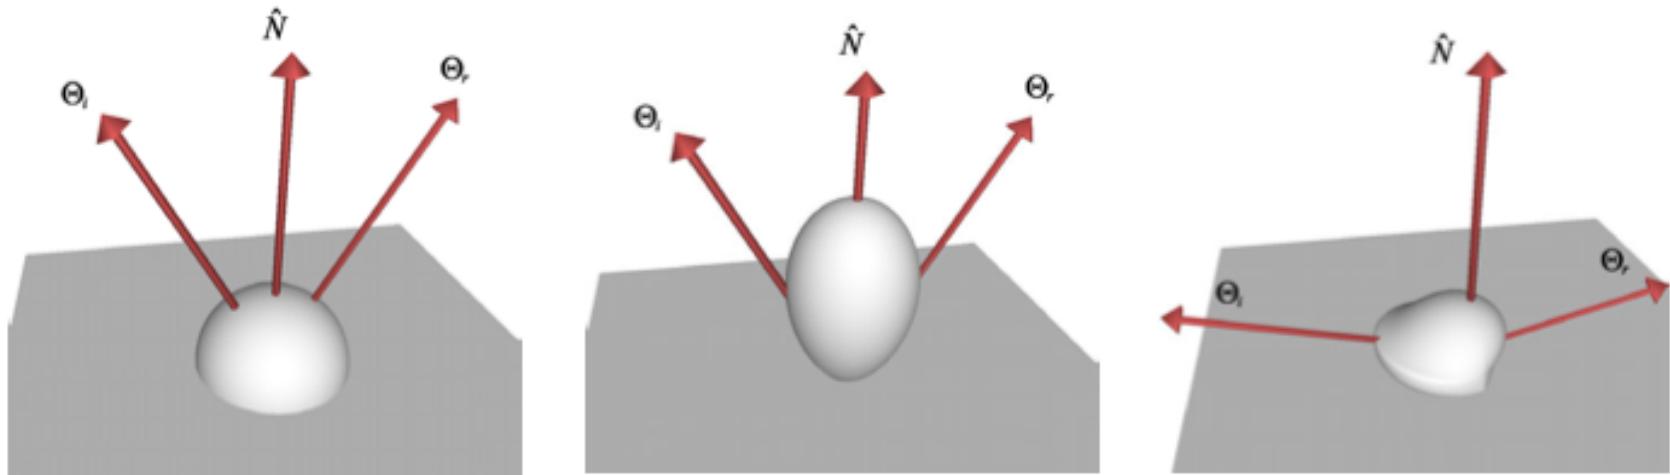
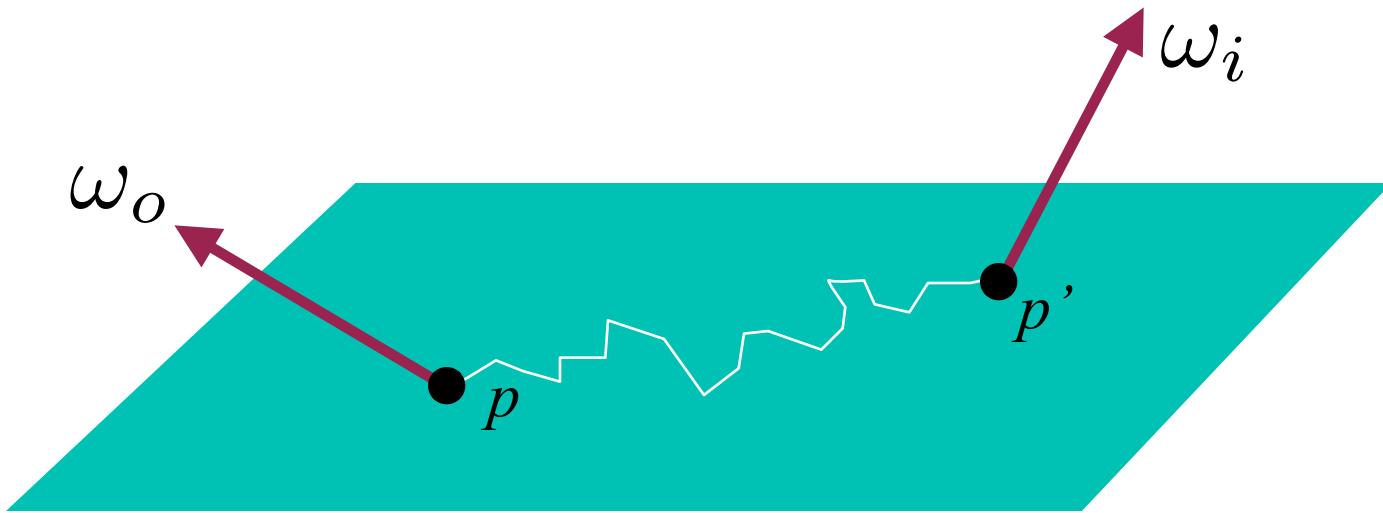


Fig. 2. Diffuse BRDF models: Lambert, Minnaert and Oren-Nayar.

BSSRF

- Bidirectional scattering-surface distribution function $S()$



- We now need to integrate over surface area as well

$$L(p, \omega_0) = \int_A \int_{S^2} S(p', \omega_i, p, \omega_o) \cos \theta_i d\omega_i dA$$

- Examples: <http://graphics.ucsd.edu/~henrik/images/subsurf.html>

Rendering Equation

Simple but difficult to solve

- High dimensionality
 - *and we have not even used a variable for color (wavelength of light)*

Solutions based on

- simplifying assumptions
- stochastic sampling
- FEM

One solution

Ray-tracing

Open source ray tracers

- Povray: www.povray.org
- LuxRender, <http://www.luxrender.net>
- MegaPov

Raytracing

Offline high quality rendering

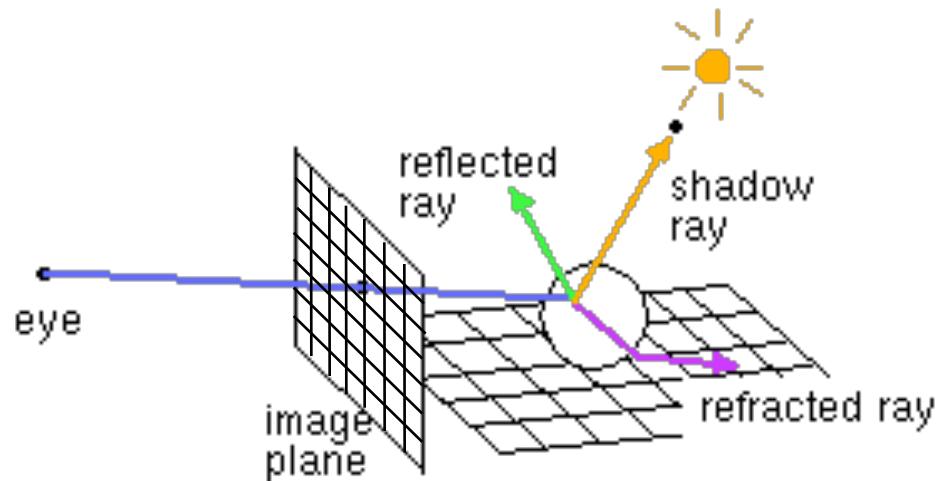
- Real-time versions for limited complexity scenes

Tuned for specular and transparent objects

- Partly physics, optics

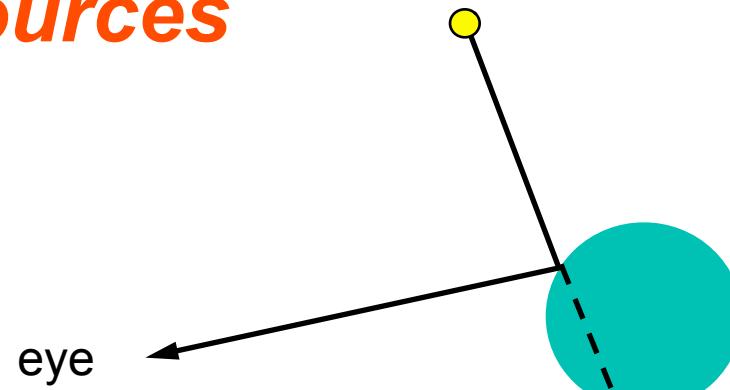
Main idea:

- A pixel should have the color of the object point that projects to it.

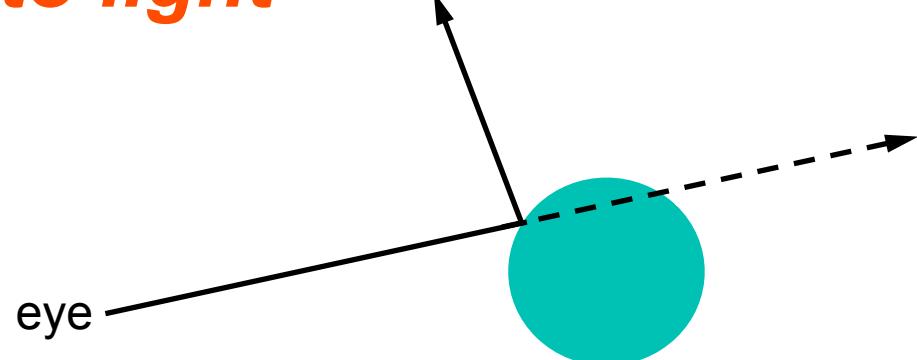


Forward and Backward methods

*Forward: from light sources
to eye*



*Backward: from eye to light
sources*



Our version

There are different variants of the backwards RT method that make different assumptions and simplifications

We will study a specific version that is

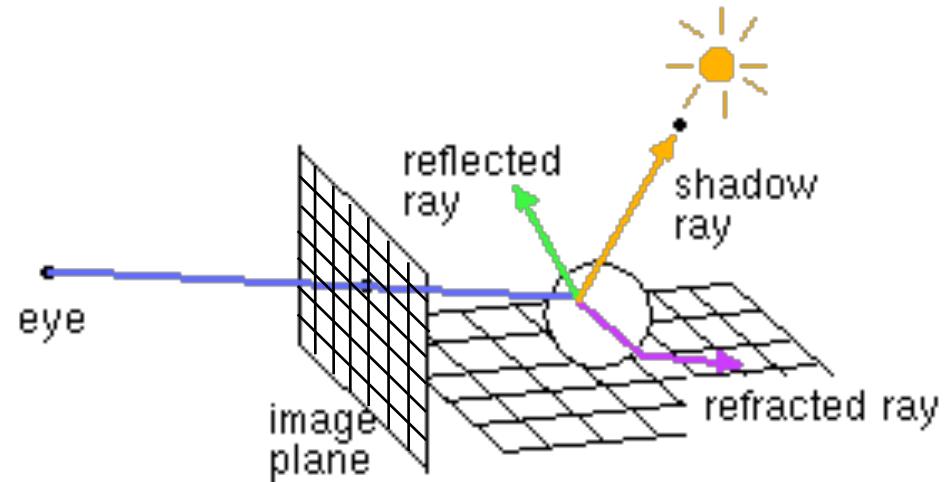
- Efficient
- Simple to implement
- Educational

Original Backwards Raytracing

Turner Whitted 1979

Assumptions

- Reflection and Refraction only on ideal direction cones
- Lambertian diffuse surfaces (no custom BRDFs)
- Diffuse objects do not receive reflected light



for each pixel on screen

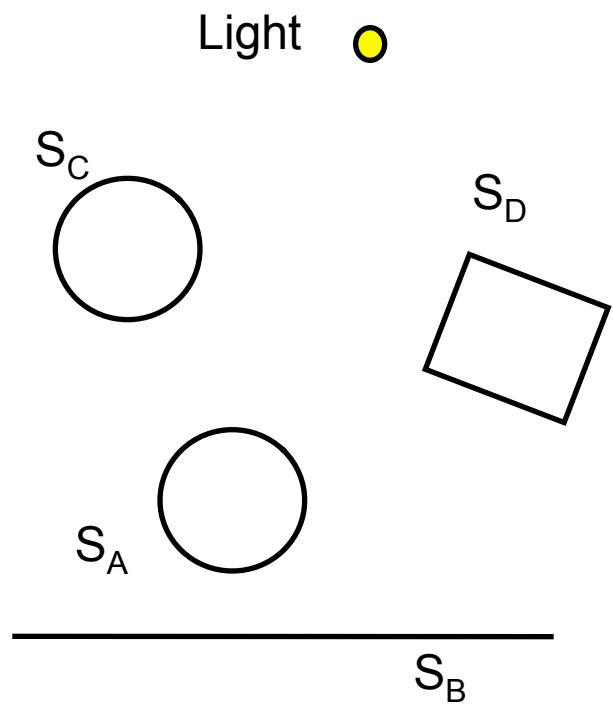
- determine ray from eye through pixel
- find closest intersection of ray with an object
- compute direct illumination (shadow ray)
- trace reflected and refracted ray, recursively
- calculate pixel colour, draw pixel

end

Scene

S_A	shiny, transparent
S_B, S_D	diffuse, opaque
S_C	shiny, opaque

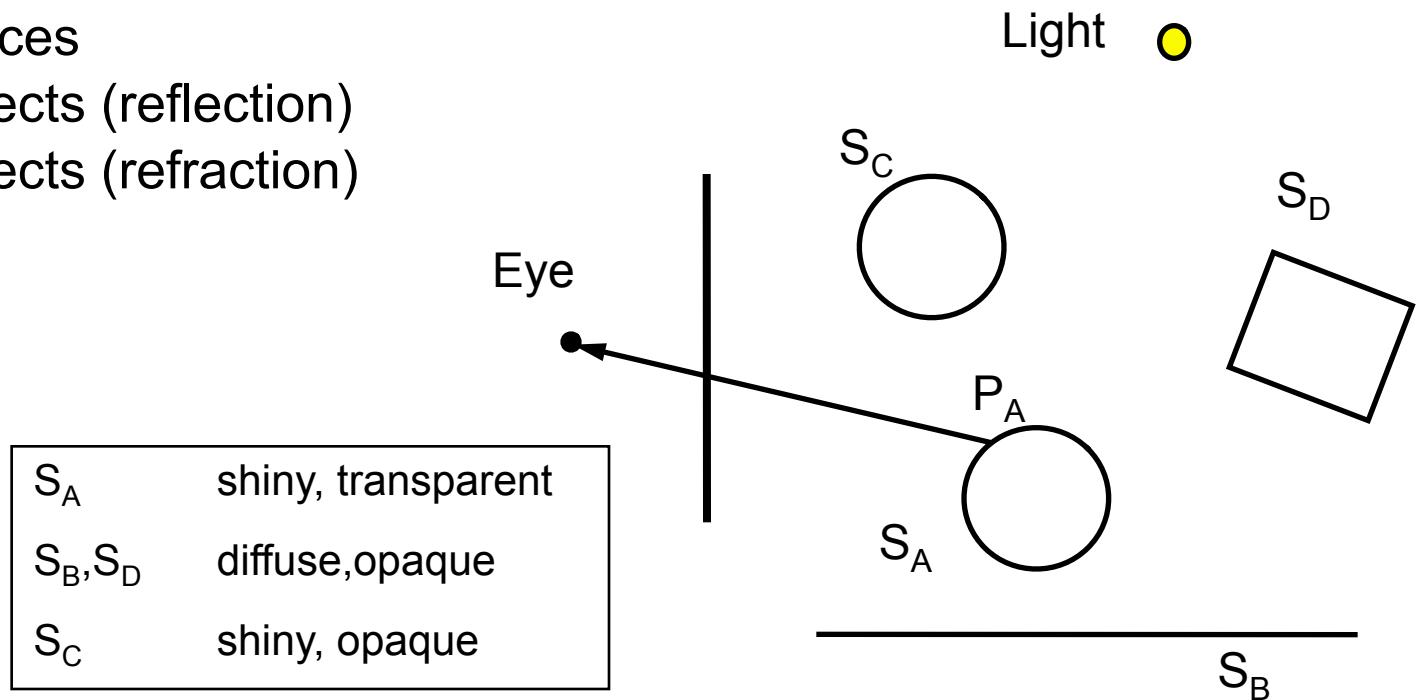
Eye



Three sources of light

The light that point P_A emits to the eye comes from:

- light sources
- other objects (reflection)
- other objects (refraction)

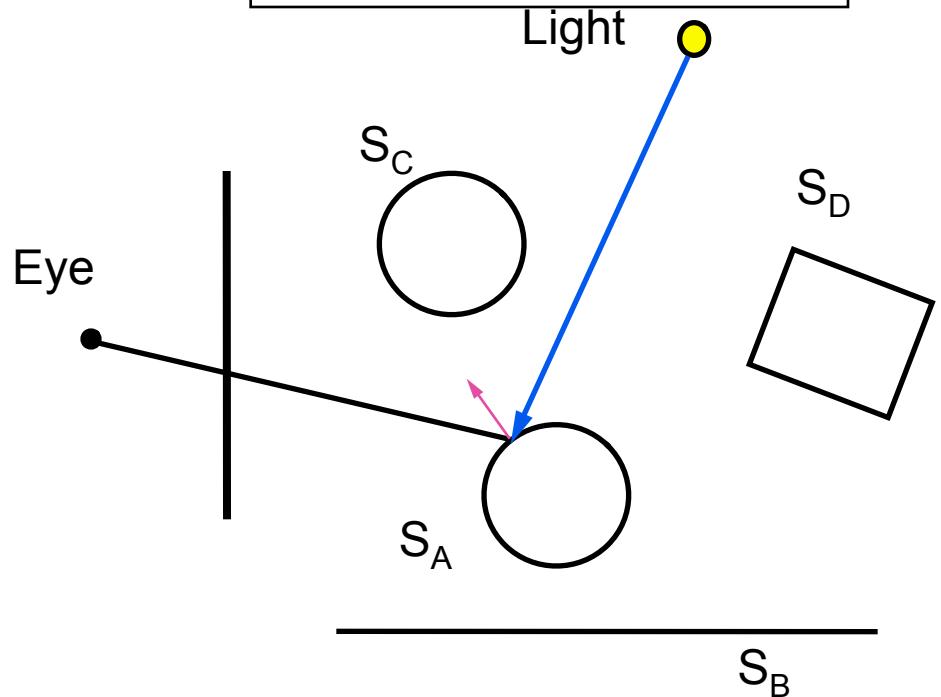


Directly from a light source

Local illumination model:

$$I = I_a + I_{diff} + I_{spec}$$

S_A	shiny, transparent
S_B, S_D	diffuse, opaque
S_C	shiny, opaque



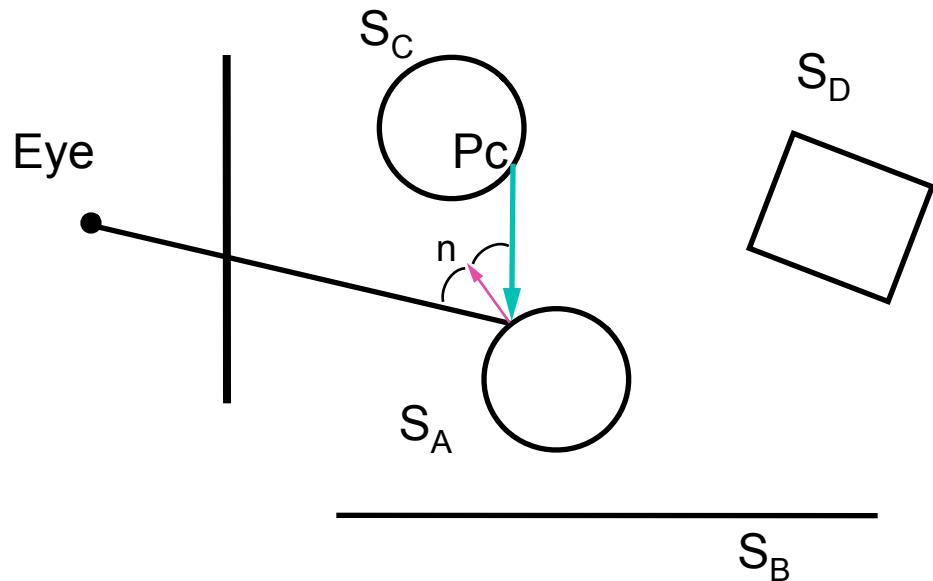
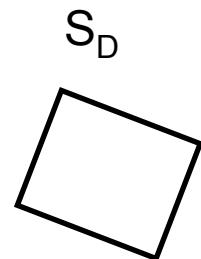
Reflection

What is the color that is reflected to P_A and that P_A reflects back to the eye?

What ever comes from the direction that eventually reflects towards the eye.

S_A	shiny, transparent
S_B, S_D	diffuse, opaque
S_C	shiny, opaque

Light



Reflection

What is the color that is reflected to P_A and that P_A reflects back to the eye?

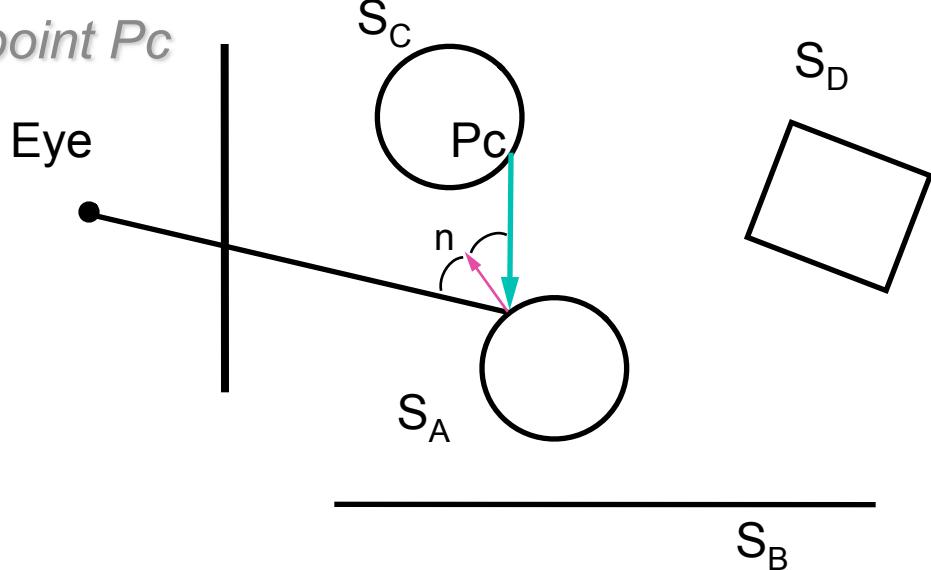
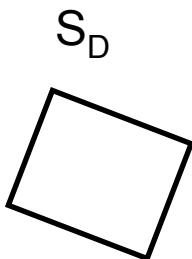
What ever comes from the direction that eventually reflects towards the eye.

That direction gets light from point P_C

What is the color of P_C ?

S_A	shiny, transparent
S_B, S_D	diffuse, opaque
S_C	shiny, opaque

Light



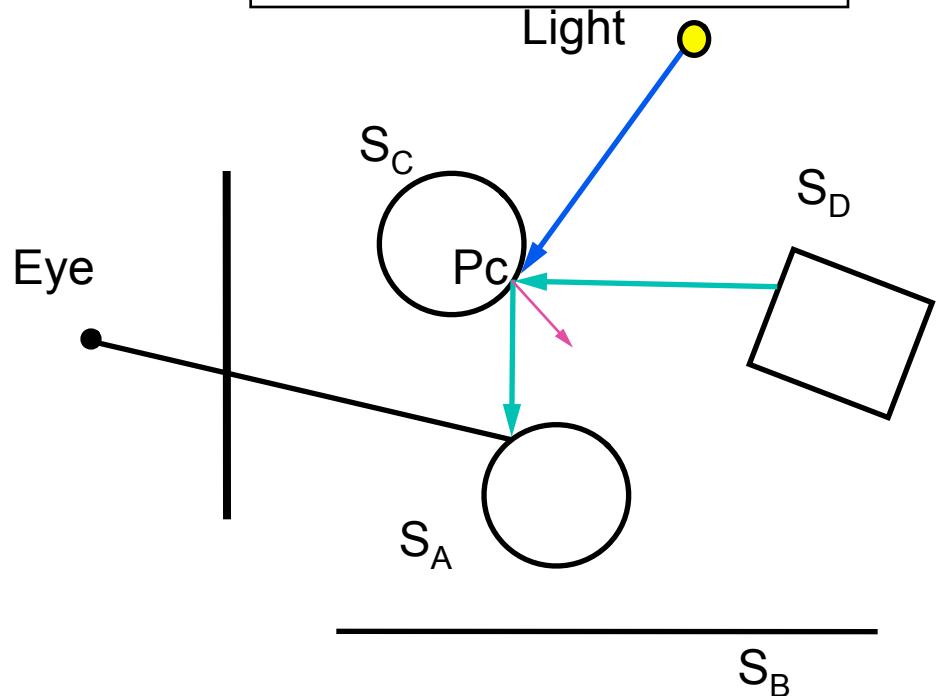
Reflection

What is the color of P_c ?

Just like P_A : raytrace P_c i.e. compute
the three contributions from

- *Light sources*
- *Reflection*
- *Refraction*

S_A	shiny, transparent
S_B, S_D	diffuse, opaque
S_C	shiny, opaque



Refraction

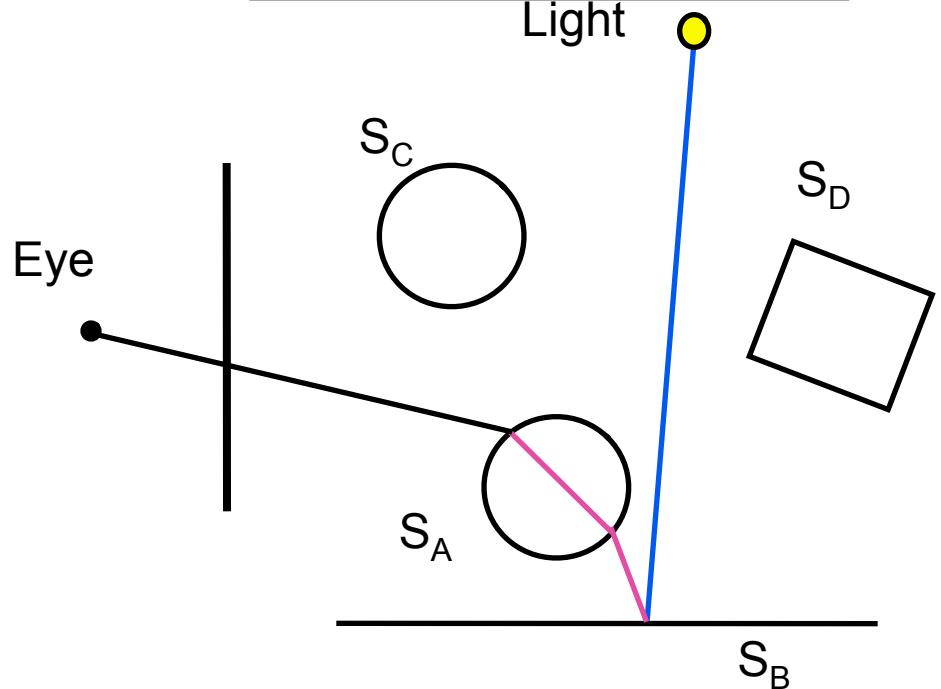
Transparent materials

How do we compute the refracted contribution?

We raytrace the refracted ray.

1. *Lights*
2. *Reflection*
3. *Refraction*

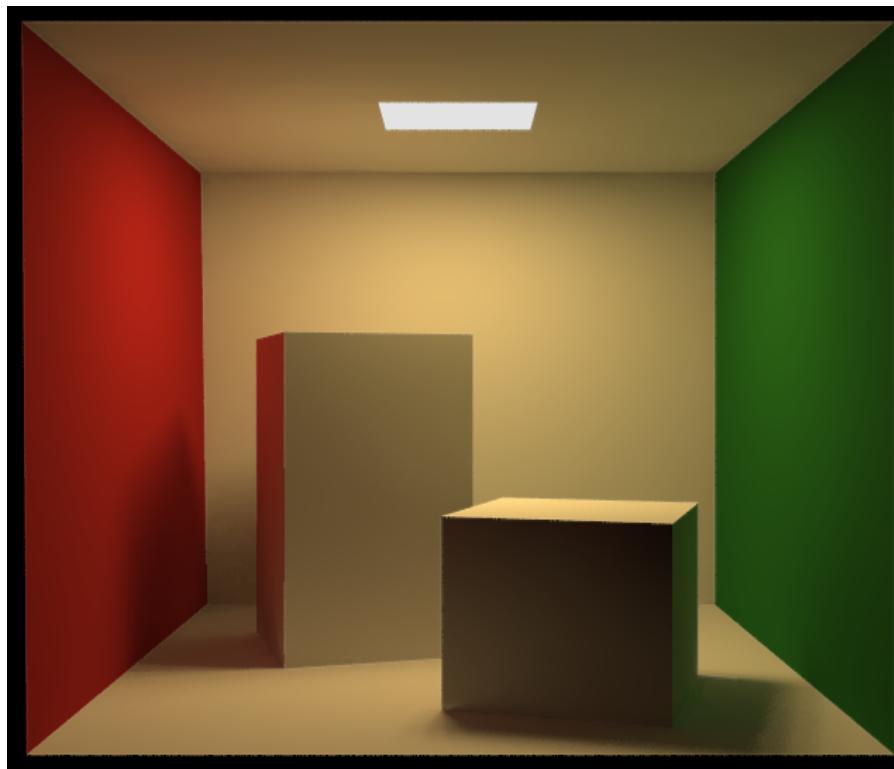
S_A	shiny, transparent
S_B, S_D	diffuse, opaque
S_C	shiny, opaque



What have we ignored ?

What have we ignored ?

Diffuse objects do not receive light from other objects.



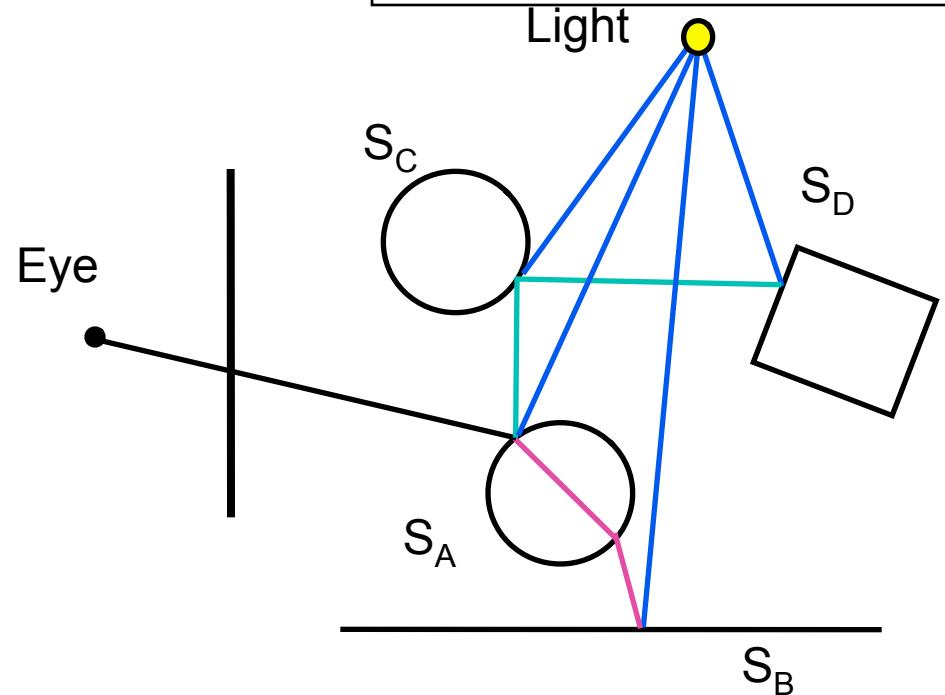
Three sources of light together

The color that the pixel is assigned comes from:

- light sources
- other objects (reflection)
- other objects (refraction)

It is more convenient to trace the rays from the eye to the scene (backwards)

S_A	shiny, transparent
S_B, S_D	diffuse, opaque
S_C	shiny, opaque



Backwards Raytracing Algorithm

For each pixel construct a ray: eye \rightarrow pixel

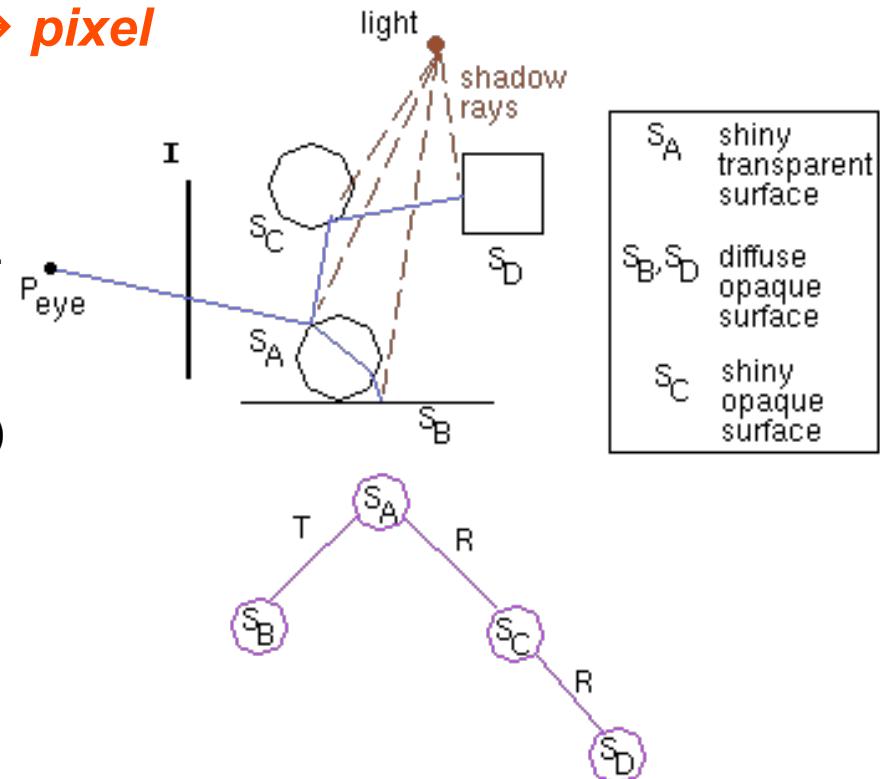
```
raytrace( ray )
```

```
P = compute_closest_intersection(ray)  
color_local = ShadowRay(light1, P)+...  
+ ShadowRay(lightN, P)
```

```
color_reflect = raytrace(reflected_ray )  
color_refract = raytrace(refracted_ray )
```

```
color = color_local  
+ kre*color_reflect  
+ kra*color_refract
```

```
return( color )
```



How many levels of recursion do we use?

The more the better.

Infinite reflections at the limit.

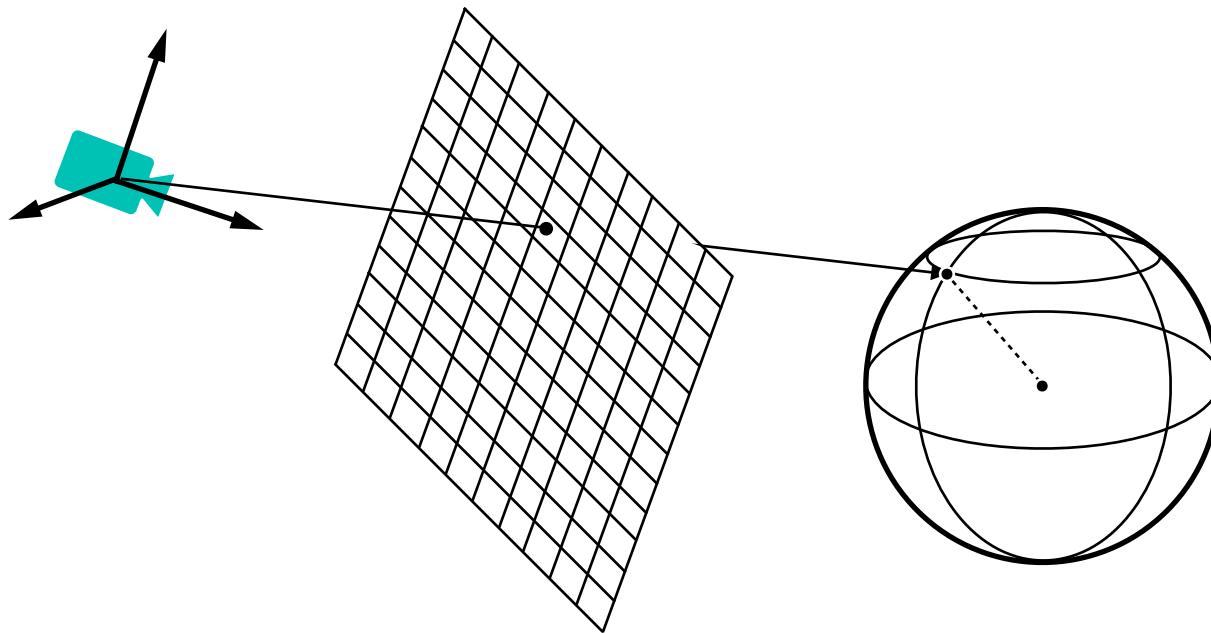


Stages of raytracing

We need to do the following:

- Set the camera and the image plane
- Compute a ray from the eye to every pixel and trace it in the scene
- Compute object-ray intersections
- Cast shadow, reflected and refracted rays at each intersection

Setting up the camera



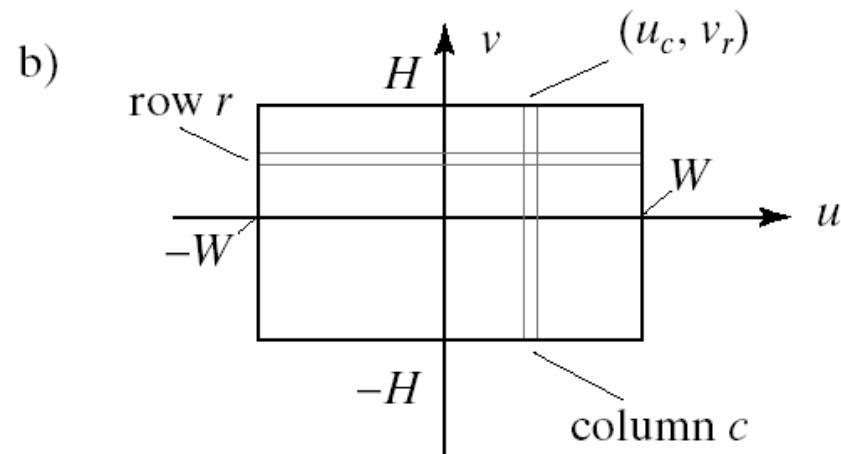
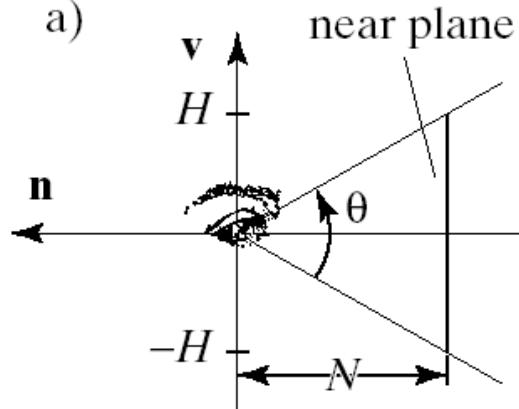
Set the viewport and resolution in Camera Coordinates

Image parameters

- Width $2W$, Height $2H$
- Number of pixels $n_{\text{Cols}} \times n_{\text{Rows}}$

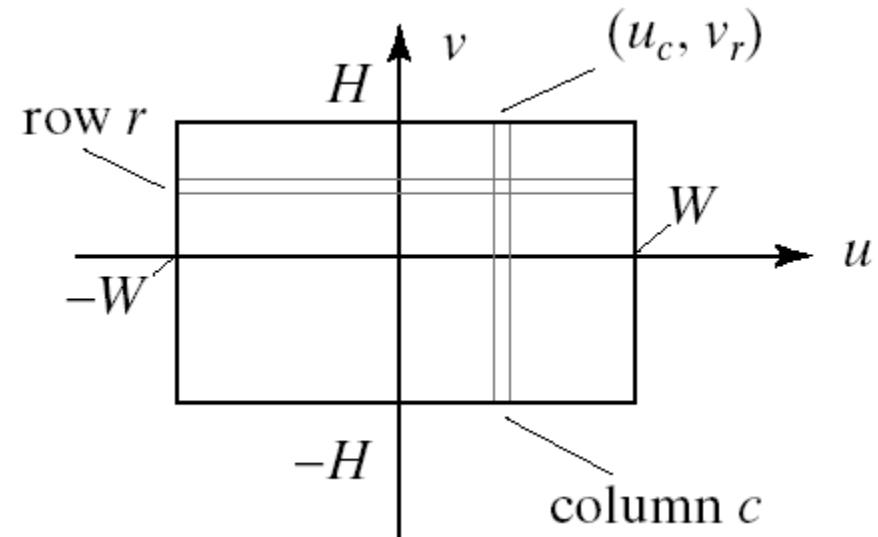
Camera parameters (view point)

- Camera coordinate system (eye , $\mathbf{u}, \mathbf{v}, \mathbf{n}$)
- Image plane at $\mathbf{n} = -N$



Pixel coordinates in camera coordinate system

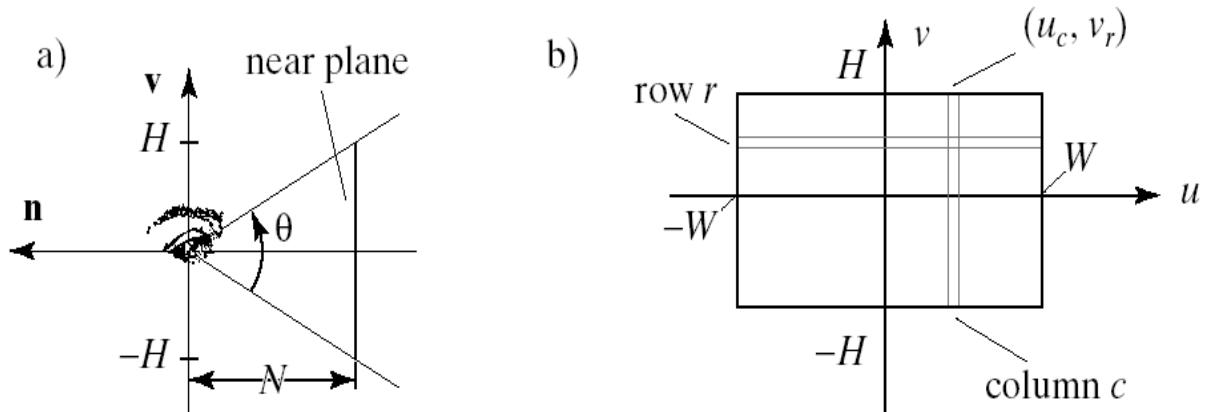
Lower left corner of pixel $P(r,c)$ has coordinates in camera space:



$$u_c = -W + W \frac{2c}{nCols}, \quad c = 0, 1, \dots, nCols - 1,$$

$$v_r = -H + H \frac{2r}{nRows}, \quad r = 0, 1, \dots, nRows - 1,$$

Ray through pixel



Lower left corner

Camera coordinates : $P(r, c) = (u_c, v_r, -N)$

World coordinates : $P(r, c) = eye - N\mathbf{n} + u_c\mathbf{u} + u_r\mathbf{v}$

Ray through pixel:

$$ray(r, c, t) = eye + t(P(r, c) - eye)$$

$$ray(r, c, t) = eye + t \left(-N\mathbf{n} + W\left(\frac{2c}{nCols} - 1\right)\mathbf{u} + H\left(\frac{2r}{nRows} - 1\right)\mathbf{v} \right)$$

Ray-object intersections

Objects

- Spheres, cylinders, surfaces, etc
- Any object that we can mathematically compute intersections with lines

For us

- Only affine transformed spheres

Sphere

Implicit

$$f(x, y, z) = (x - x_0)^2 + (y - y_0)^2 + (z - z_0)^2 = R^2$$

Parametric

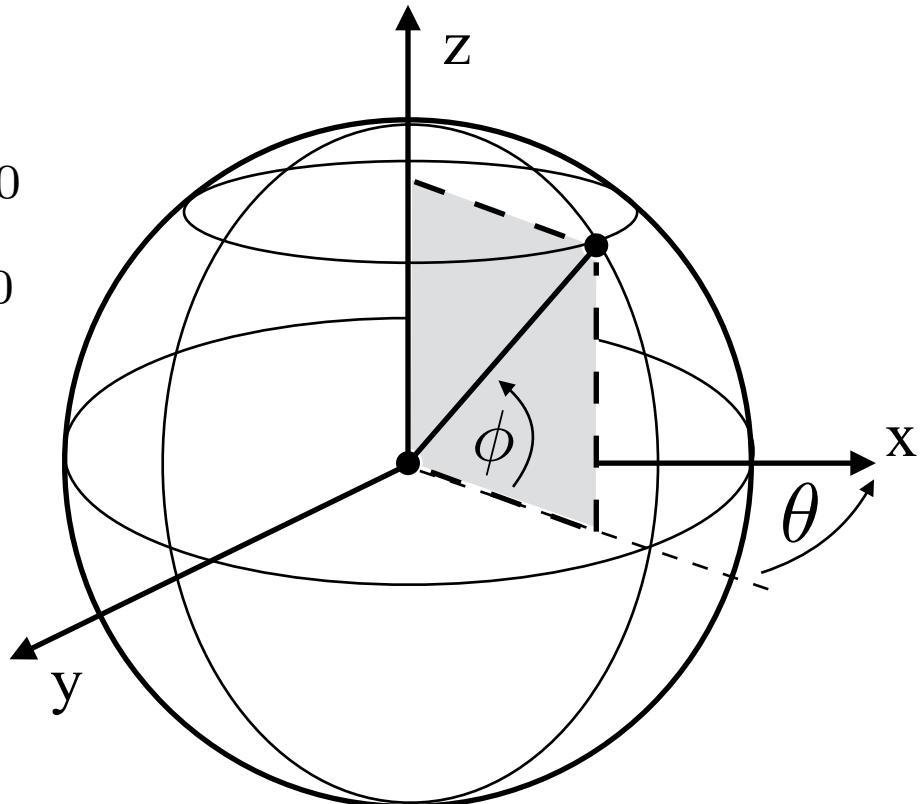
$$x(\phi, \theta) = R\cos(\phi)\cos(\theta) + x_0$$

$$y(\phi, \theta) = R\cos(\phi)\sin(\theta) + y_0$$

$$z(\phi, \theta) = R\sin(\phi) + z_0$$

$$-\pi/2 \leq \phi \leq \pi/2$$

$$-\pi \leq \theta \leq \pi$$



Canonical Sphere

- Radius 1 centered around the origin
- Equivalent forms

$$\sqrt{x^2 + y^2 + z^2} = 1$$

$$x^2 + y^2 + z^2 = 1$$

$$|P| - 1 = 0$$

$$|P| = 1 \rightarrow \sqrt{P \cdot P} = 1 \rightarrow P \cdot P = 1$$

Normal on a sphere

Normal

$$\mathbf{N}(x, y, z) = \nabla f(x, y, z) = 2(x - x_o, y - y_0, z - z_0)$$

Unit Normal

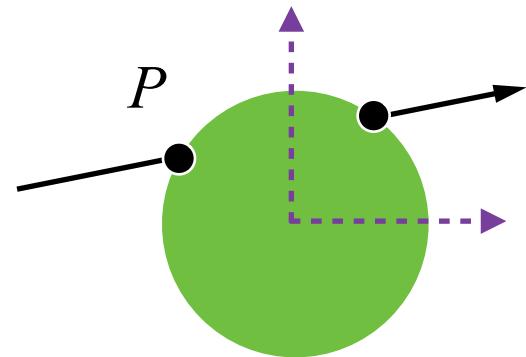
$$\mathbf{n}(x, y, z) = \mathbf{N}/\|\mathbf{N}\|$$

Ray - Canonical sphere intersection

Can I use homogeneous coordinates?

$$ray(t) = S + \mathbf{c}t$$

$$Sphere(P) = |P| - 1 = 0$$



$$Sphere(ray(t)) = 0 \Rightarrow$$

$$|S + \mathbf{c}t| - 1 = 0 \Rightarrow (S + \mathbf{c}t)(S + \mathbf{c}t) - 1 = 0 \Rightarrow$$

$$|\mathbf{c}|^2 t^2 + 2(S \cdot \mathbf{c})t + |S|^2 - 1 = 0$$

This is a quadratic equation in t

Solving a quadratic equation

$$|\mathbf{c}|^2 t^2 + 2(S \cdot \mathbf{c})t + |S|^2 - 1 = 0$$
$$At^2 + 2Bt + C = 0$$

$$t_h = -\frac{B}{A} \pm \frac{\sqrt{B^2 - AC}}{A}$$
$$t_h = -\frac{S \cdot \mathbf{c}}{|\mathbf{c}|^2} \pm \frac{\sqrt{(S \cdot \mathbf{c})^2 - |\mathbf{c}|^2 (|S|^2 - 1)}}{|\mathbf{c}|^2}$$

If $(B^2 - AC) = 0$ one solution

If $(B^2 - AC) < 0$ no solution

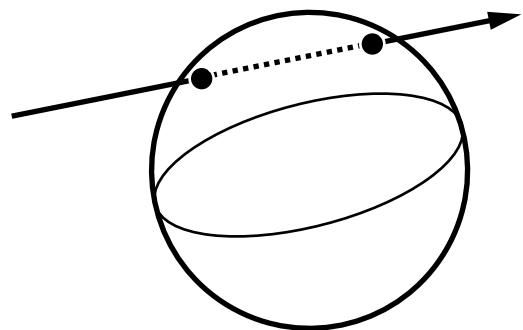
If $(B^2 - AC) > 0$ two solutions

Note!

Can I use homogeneous coordinates?

$$ray(t) = S + ct$$

$$Sphere(P) = |P| - 1 = 0$$



$$Sphere(ray(t)) = 0 \Rightarrow$$

$$|S + ct| - 1 = 0 \Rightarrow (S + ct)(S + ct) - 1 = 0 \Rightarrow$$

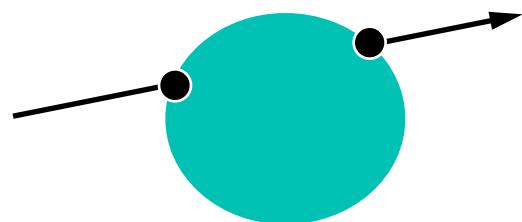
$$|\mathbf{c}|^2 t^2 + 2(S \cdot \mathbf{c})t + |S|^2 - 1 = 0$$

Note!

Can I use homogeneous coordinates?

$$ray(t) = S + ct$$

$$Sphere(P) = |P| - 1 = 0$$



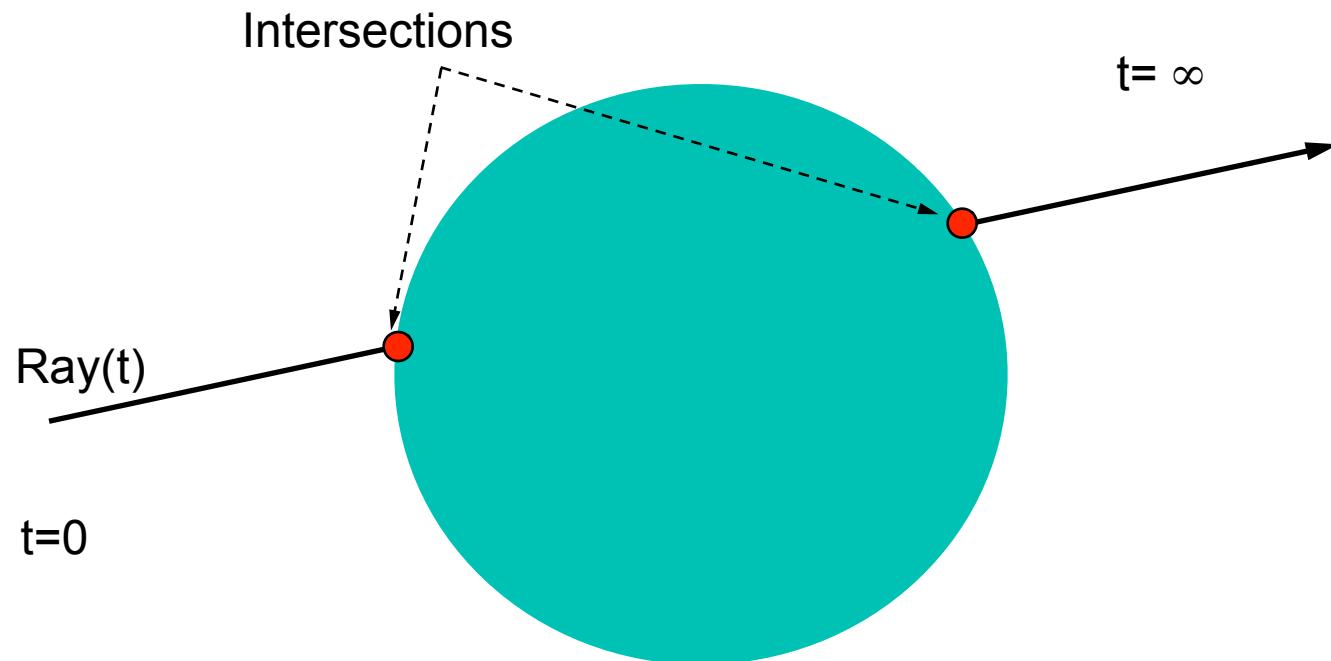
$$Sphere(ray(t)) = 0 \Rightarrow$$

$$|S + ct| - 1 = 0 \Rightarrow (S + ct)(S + ct) - 1 = 0 \Rightarrow$$

$$|\mathbf{c}|^2 t^2 + 2(S \cdot \mathbf{c})t + |S|^2 - 1 = 0$$

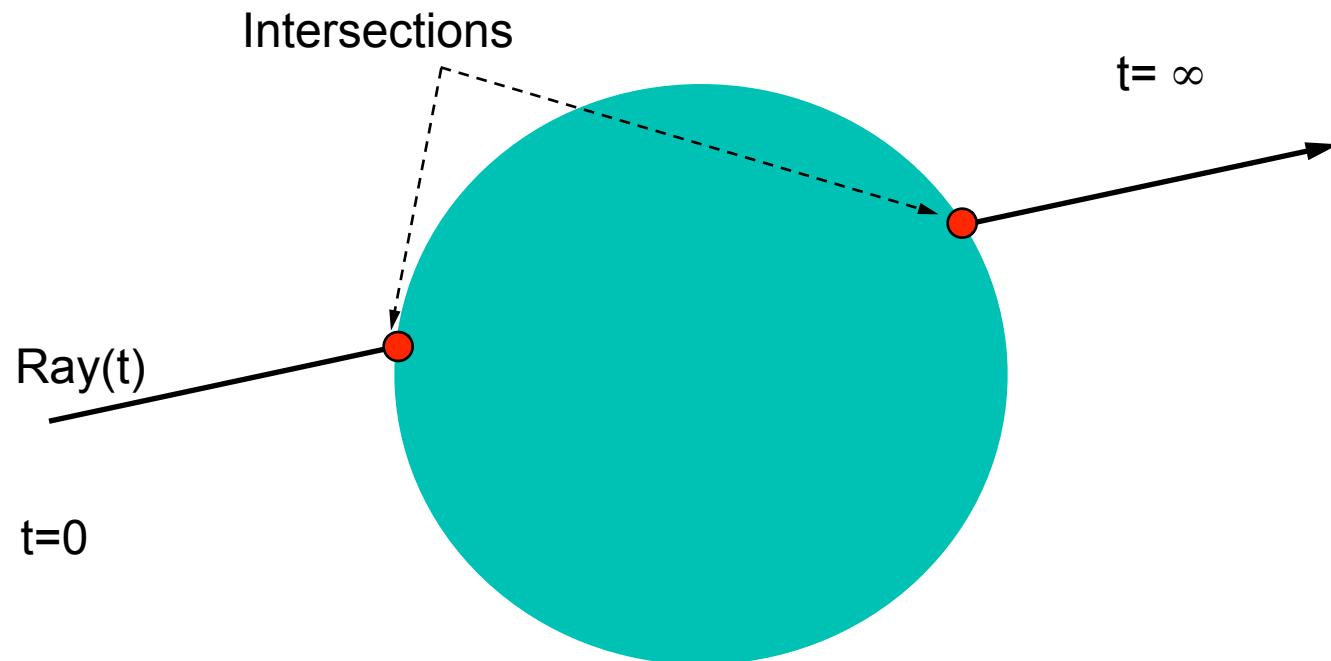
Must solve the equation in cartesian space

First intersection?



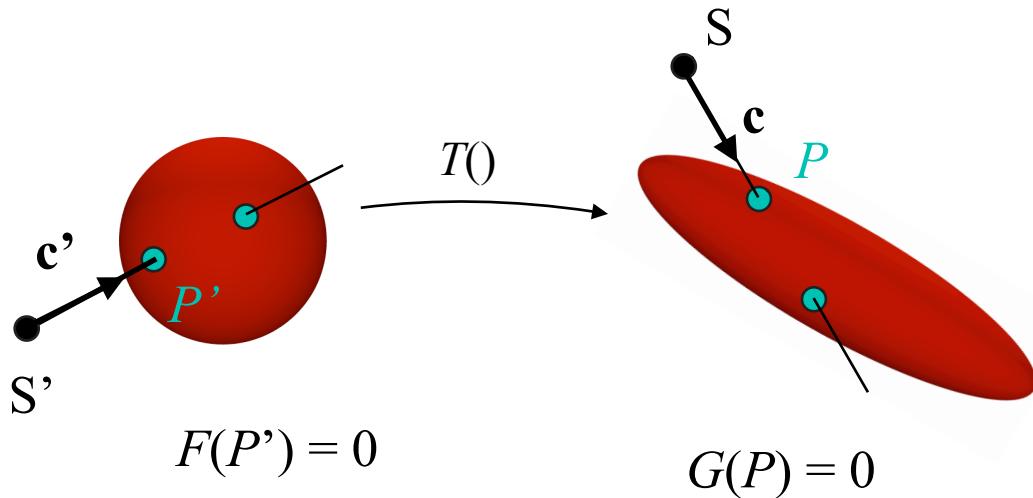
First intersection?

$t_1 < t_2$

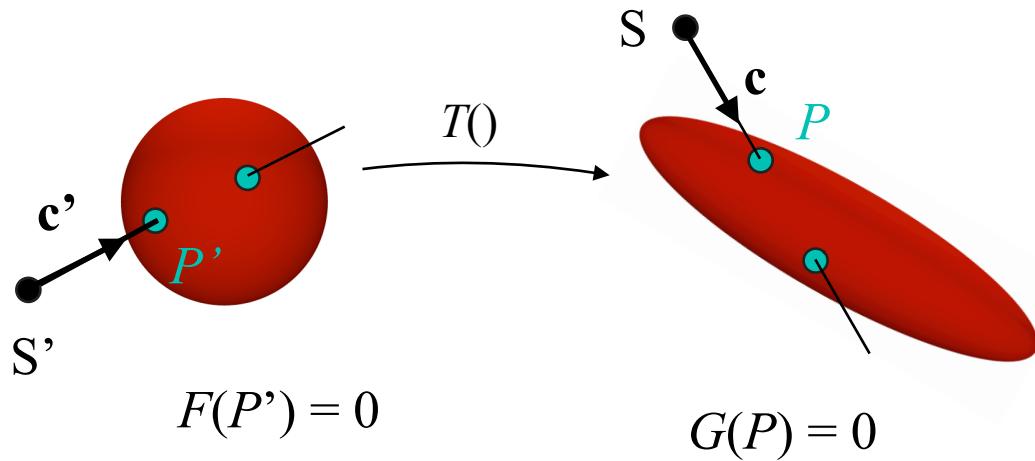


Transformed primitives?

- That was a canonical sphere.
- Where does $S+ct$ hit the transformed sphere $G = T(F)$?



Linear transformation

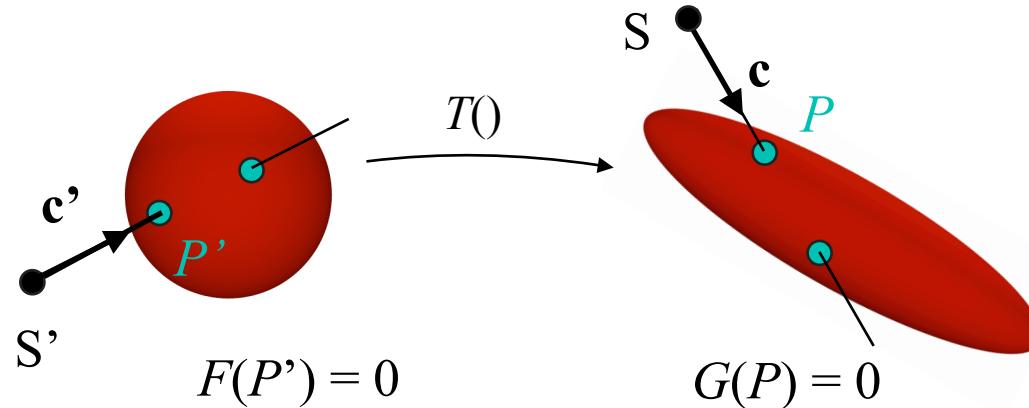


Implicit equation $G(P) = 0$.

Untransformed implicit equation $F(P') = 0$.

$$P = MP' \Rightarrow P' = M^{-1}P$$

Linear transformation



$$P = MP' \Rightarrow P' = M^{-1}P$$

$$F(P') = F(T^{-1}(P)) = 0 \Rightarrow F(T^{-1}(P)) = 0$$

$$F(T^{-1}(S + ct)) = 0 \Rightarrow$$

$$F(T^{-1}(S) + T^{-1}(ct)) = 0$$

Which means that we can intersect the inverse transformed ray with the untransformed primitive.

Final Intersection

Inverse transformed ray

$$\tilde{r}(t) = M^{-1} \begin{pmatrix} S_x \\ S_y \\ S_z \\ 1 \end{pmatrix} + t M^{-1} \begin{pmatrix} c_x \\ c_y \\ c_z \\ 0 \end{pmatrix} = \tilde{S}' + \tilde{c}'t$$

- Drop 1 and 0 to get $S' + c't$.

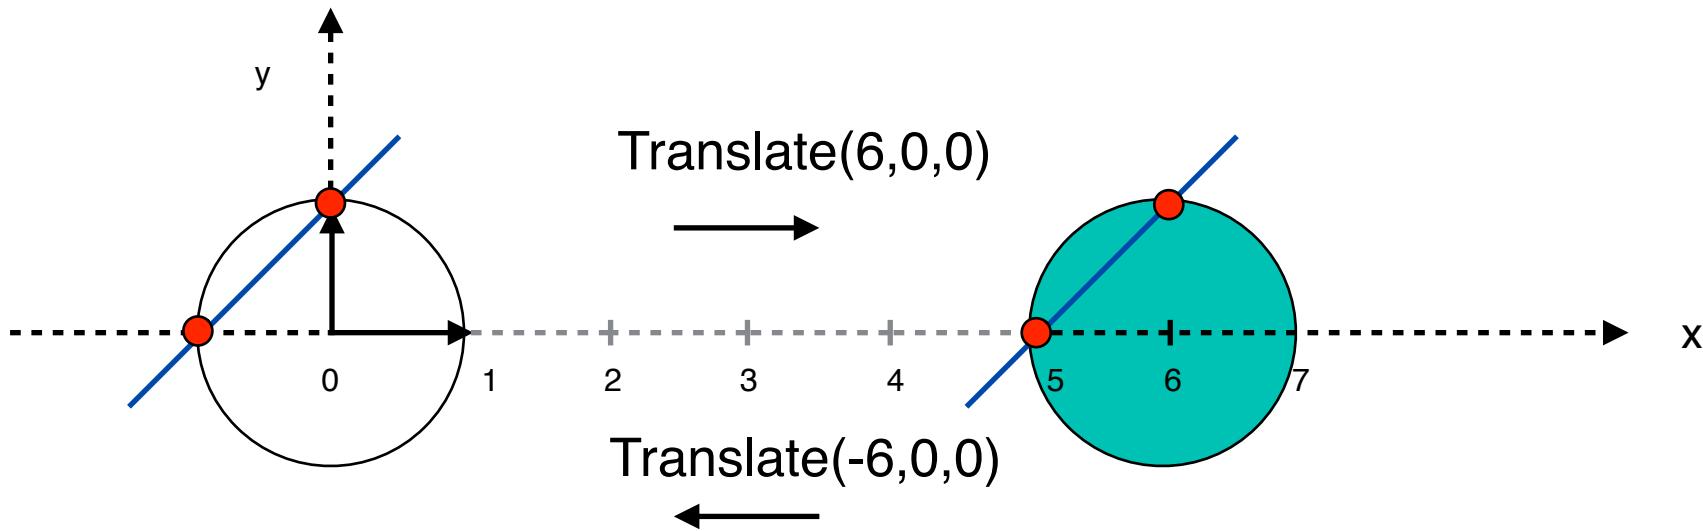
So ..for each object

- Inverse transform ray and get $S' + c't$.
- Find the intersection t, t_h , between inv-ray and canonical object.
- Use t_h in the untransformed ray $S + ct$ to find the intersection.
- **For the normal, use the inverse transpose**

Example

Canonical sphere, radius 1
 $|x,y,z| - 1 = 0$

Translated sphere
 $|x-6,y,z| - 1 = 0$



$$\text{Ray}'(t) = (-1,0,0) + t(1,1,0)$$

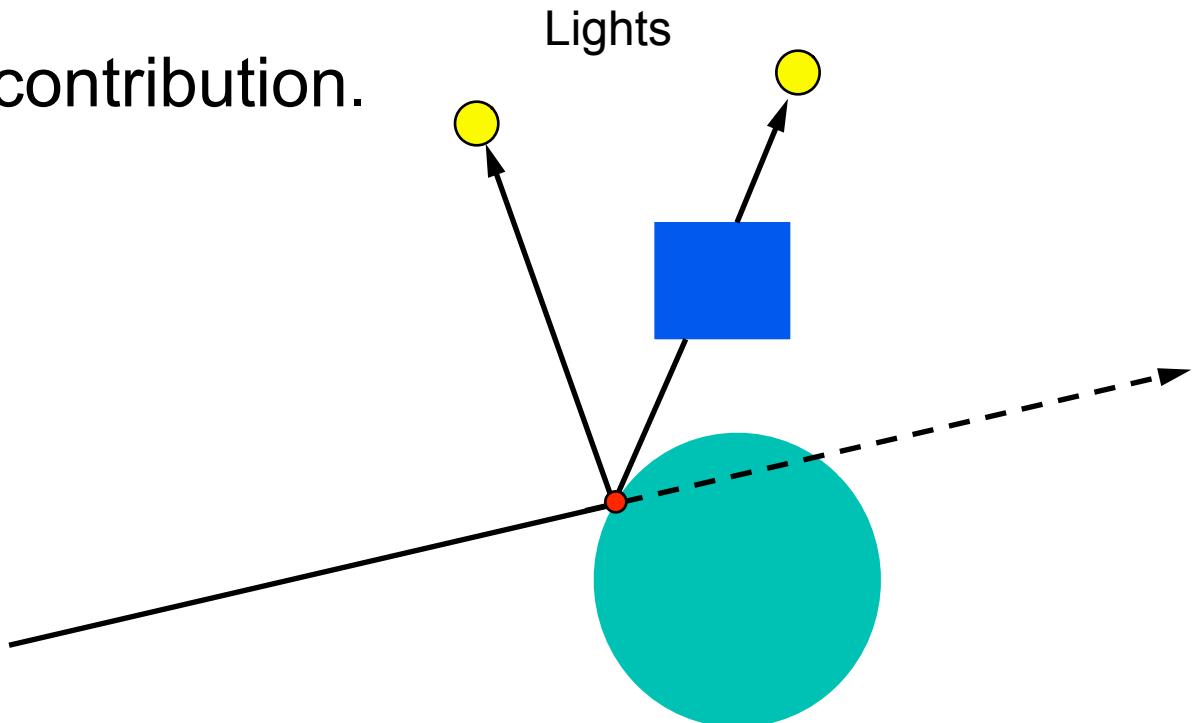
Intersections
(-1,0,0) and (0,1,0)

$$\text{Ray}(t) = (5,0,0) + t(1,1,0)$$

Intersections:
(5,0,0) and (6,1,0)

Shadow ray

- For each light intersect shadow ray with all objects.
- If no intersection is found apply local illumination at intersection.
- If in shadow no contribution.



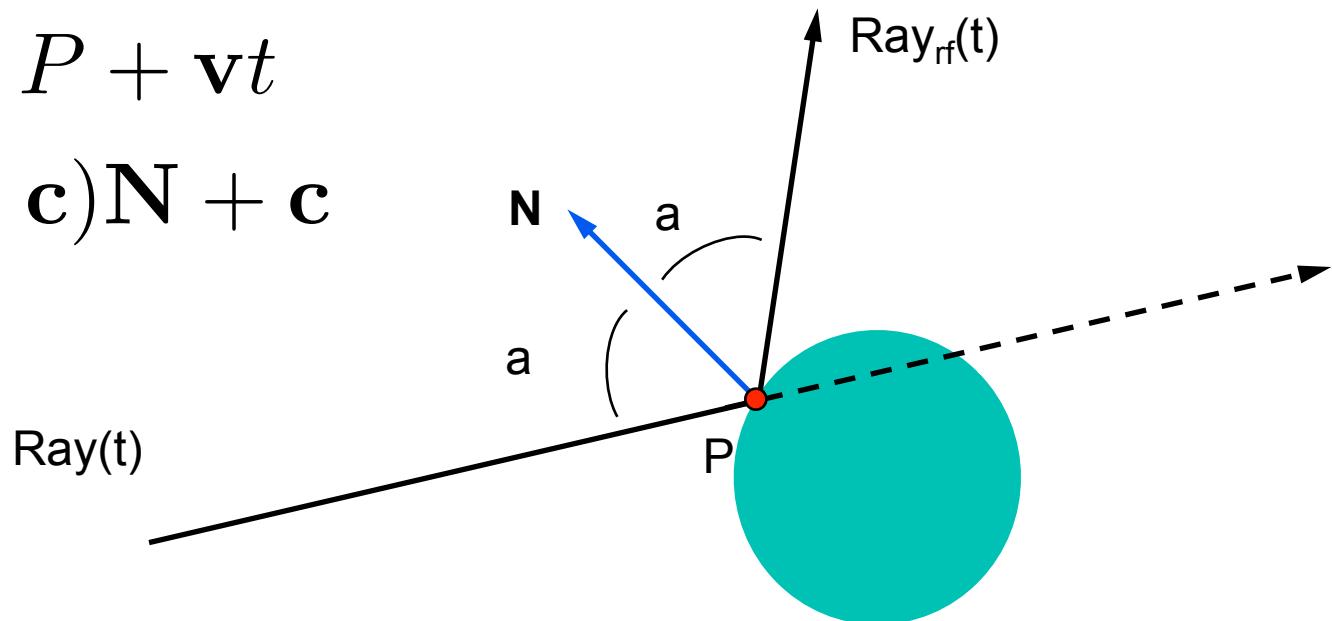
Reflected ray

Raytrace the reflected ray

$$Ray(t) = A + \mathbf{c}t$$

$$Ray_{rf}(t) = P + \mathbf{v}t$$

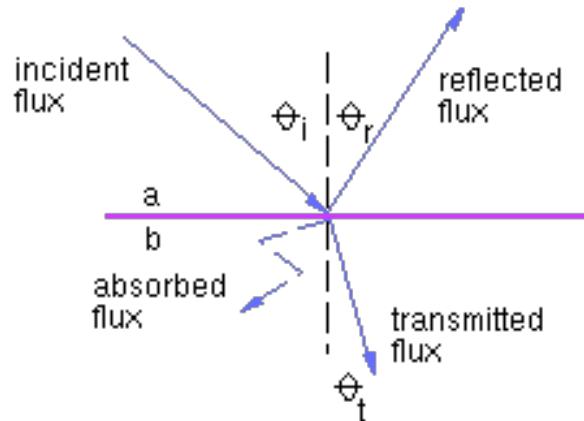
$$\mathbf{v} = -2(\mathbf{N} \cdot \mathbf{c})\mathbf{N} + \mathbf{c}$$



Refracted ray

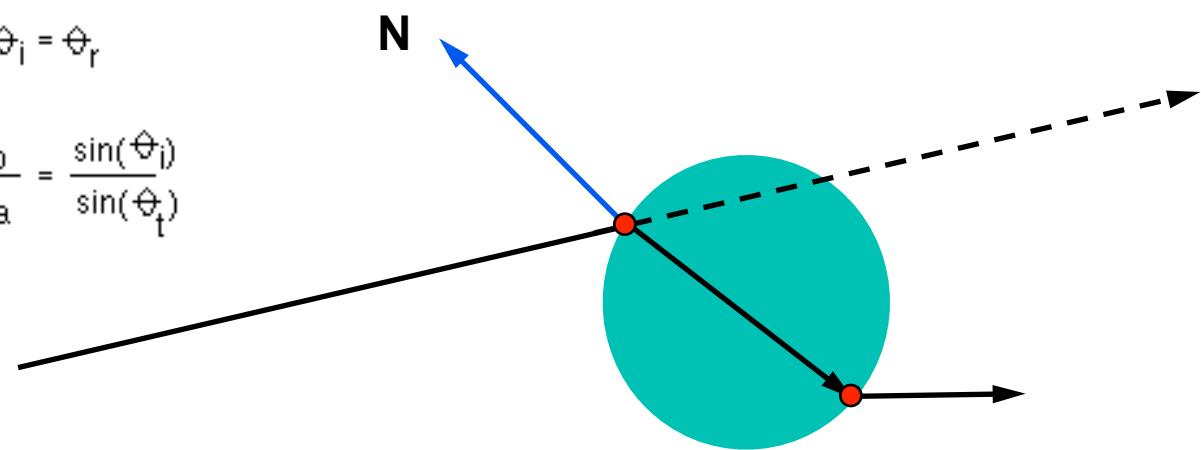
Raytrace the refracted ray

Snell's law



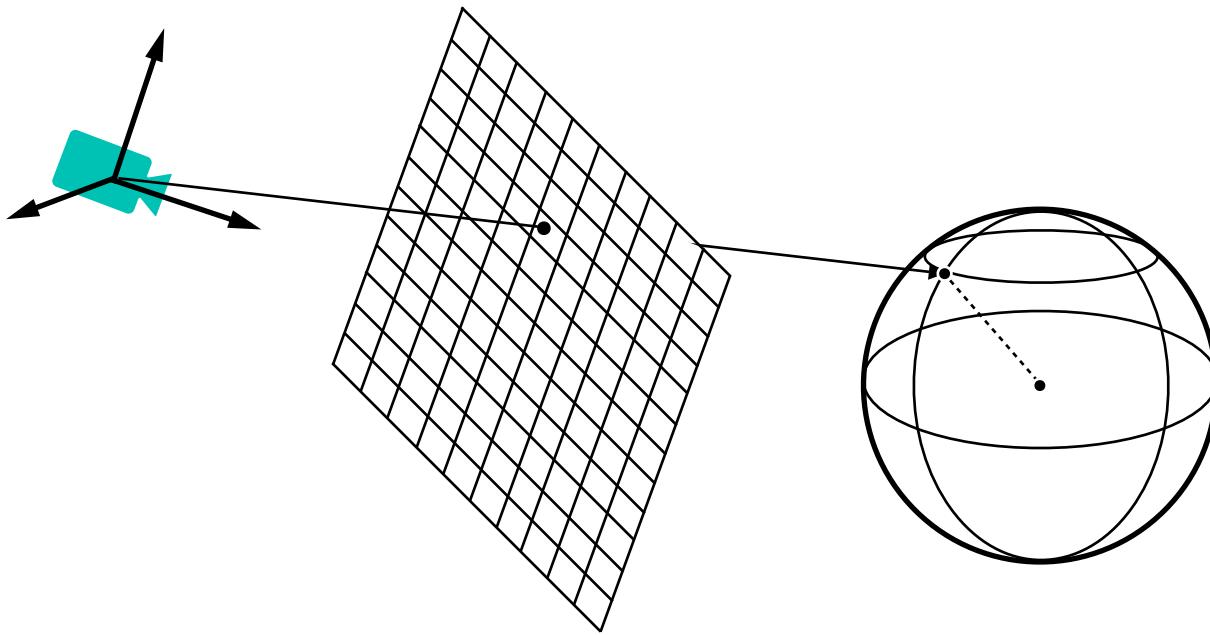
$$\theta_i = \theta_r$$

$$\frac{n_b}{n_a} = \frac{\sin(\theta_i)}{\sin(\theta_t)}$$



Add all together

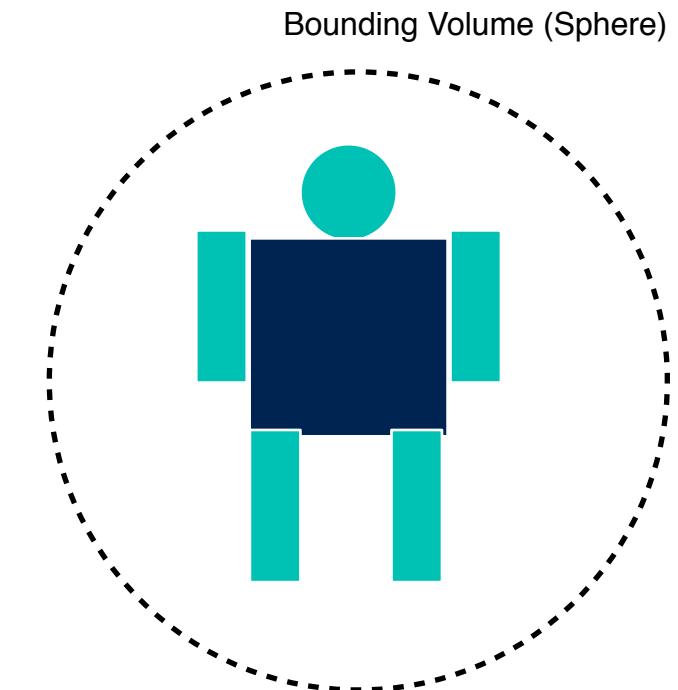
- $\text{color}(r, c) = \text{color_shadow_ray} + K_{\text{re}} * \text{color}_{\text{re}} + K_{\text{ra}} * \text{color}_{\text{ra}}$



Efficiency issues

Computationally expensive

- avoid intersection calculations
 - *Voxel grids*
 - *BSP trees*
 - *Octrees*
 - *KD-Trees*
 - *Bounding volume trees*
- optimize intersection calculations
 - *try recent hit first*
 - *reuse info from numerical methods*

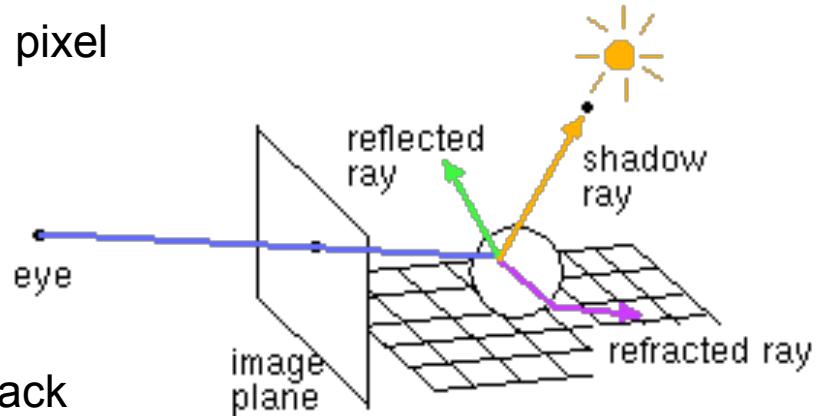


Summary: Raytracing

Recursive algorithm

```
Function Main
    for each pixel (c,r) on screen
        determine ray  $r_{c,r}$  from eye through pixel
        ray.setDepth(1)
        color(c,r) = raytrace( $r_{c,r}$ )
    end for
end

function raytrace(r)
    if (ray.depth() > MAX_DEPTH) return black
    P = closest intersection of ray with all objects
    if ((no intersection ) and (ray.depth ==1)) return backgroundColor
    clocal = Sum(shadowRays(P,Lighti))
    cre = raytrace( $r_{re}$ )
    cra = raytrace( $r_{ra}$ )
    return (clocal+kre*cre+kra*cra)
end
```



Advanced concepts

Participating media

Transculency

Sub-surface scattering (e.g. Human skin)

Photon mapping

Photon Mapping

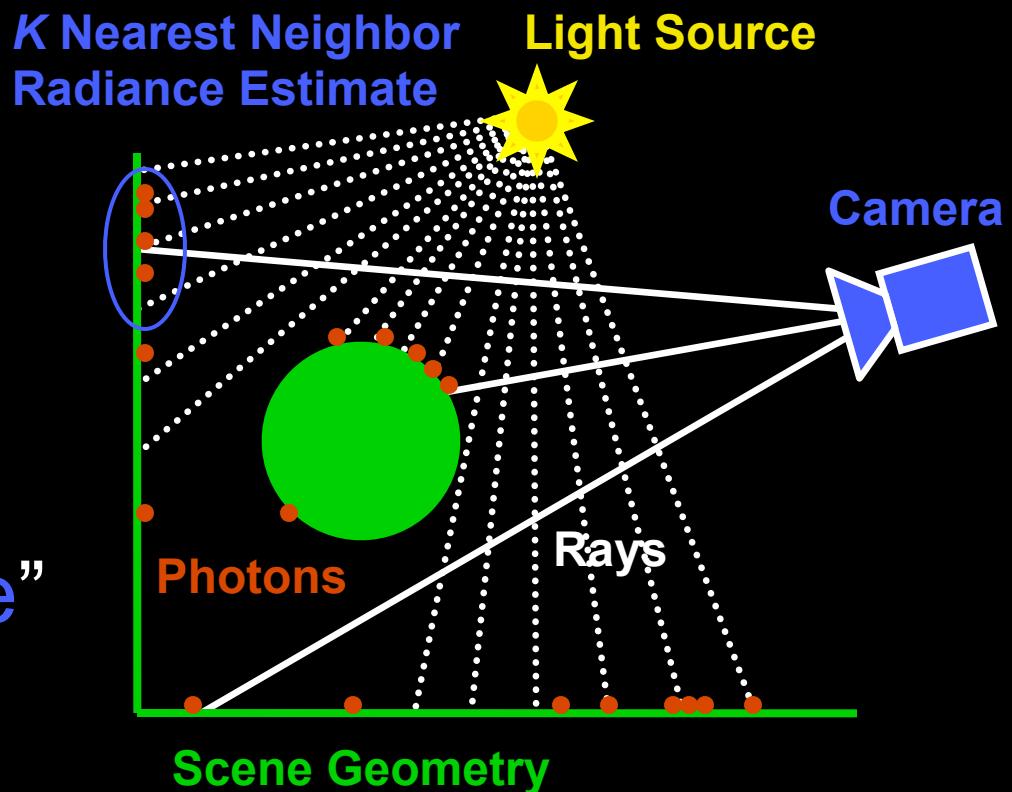
(slides courtesy of Shawn Singh)

Create a sampling of the illumination

- For scenes that do not change can be pre-computed
- View independent

Photon Mapping

- “Photons” are sample points traced from light sources
- For every ray, a “radiance estimate” is computed using photons



Resources

- Real-time Ray Tracing
Our implementation based on [Wald 2004]
- GPU Photon Mapping
[Purcell et al 2003], [Larsen and Christensen 2004]
- Specialized Hardware
[Steinhurst 2007], [Singh 2007]
- Improve performance of each k-nearest neighbor query
[Ma and McCool 2002], [Gunther et al 2004], [Wald et al. 2004]
- Improve coherence of queries
[Steinhurst et al. 2005], [Havran et al. 2005]
- Reverse photon mapping
[Havran et al 2005]

Raytracing summary

View dependent

Computationally expensive

Good for refraction and reflection effects